

July 20, 2013

Chapter 9: Nonlinear Systems and Optimization

Uri M. Ascher and Chen Greif
Department of Computer Science
The University of British Columbia
{ascher,greif}@cs.ubc.ca

Slides for the book

A First Course in Numerical Methods (published by SIAM, 2011)

<http://www.ec-securehost.com/SIAM/CS07.html>

Goals of this chapter

- To devise and assess algorithms for (continuous) unconstrained optimization, a problem setting which arises frequently in applications;
- to solve systems of nonlinear equations, thus extending Chapter 3;
- to consider more advanced conditions and techniques for constrained optimization.

Outline

- Systems of nonlinear equations
- Unconstrained optimization
- *Constrained optimization

*advanced

Systems of equations

- Equipped with knowledge on how to solve scalar nonlinear equations as well as linear systems, it is time to combine them and consider systems of nonlinear equations

$$\begin{aligned}f_1(x_1, x_2, \dots, x_n) &= 0, \\f_2(x_1, x_2, \dots, x_n) &= 0, \\&\vdots \\f_n(x_1, x_2, \dots, x_n) &= 0.\end{aligned}$$

- In vector notation, write this as $\mathbf{f}(\mathbf{x}) = \mathbf{0}$.
- Chapter 3 promises that Newton's method extends directly to the present problem. But Newton's method requires the derivative of f , so here we have to extend the concept of derivative first.

Systems of equations

- Equipped with knowledge on how to solve scalar nonlinear equations as well as linear systems, it is time to combine them and consider systems of nonlinear equations

$$\begin{aligned}f_1(x_1, x_2, \dots, x_n) &= 0, \\f_2(x_1, x_2, \dots, x_n) &= 0, \\&\vdots \\f_n(x_1, x_2, \dots, x_n) &= 0.\end{aligned}$$

- In vector notation, write this as $\mathbf{f}(\mathbf{x}) = \mathbf{0}$.
- Chapter 3 promises that Newton's method extends directly to the present problem. But Newton's method requires the derivative of f , so here we have to extend the concept of derivative first.

Multivariate Taylor expansion

- Let $\mathbf{x} = \begin{pmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{pmatrix}$, function $\mathbf{f} = \begin{pmatrix} f_1 \\ f_2 \\ \vdots \\ f_m \end{pmatrix}$, direction $\mathbf{p} = \begin{pmatrix} p_1 \\ p_2 \\ \vdots \\ p_n \end{pmatrix}$.

Assume that \mathbf{f} is sufficiently smooth (at least two bounded derivatives).

- Then Taylor expansion gives

$$\mathbf{f}(\mathbf{x} + \mathbf{p}) = \mathbf{f}(\mathbf{x}) + J(\mathbf{x})\mathbf{p} + \mathcal{O}(\|\mathbf{p}\|^2),$$

where $J(\mathbf{x})$ is the **Jacobian matrix** of first derivatives of \mathbf{f} at \mathbf{x} ,

$$J(\mathbf{x}) = \begin{pmatrix} \frac{\partial f_1}{\partial x_1} & \frac{\partial f_1}{\partial x_2} & \cdots & \frac{\partial f_1}{\partial x_n} \\ \frac{\partial f_2}{\partial x_1} & \frac{\partial f_2}{\partial x_2} & \cdots & \frac{\partial f_2}{\partial x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial f_m}{\partial x_1} & \frac{\partial f_m}{\partial x_2} & \cdots & \frac{\partial f_m}{\partial x_n} \end{pmatrix}.$$

Newton's method

- **Derivation:** By Taylor series,

$$\mathbf{f}(\mathbf{x}) = \mathbf{f}(\mathbf{x}_k) + J(\mathbf{x}_k)(\mathbf{x} - \mathbf{x}_k) + \mathcal{O}(\|\mathbf{x} - \mathbf{x}_k\|^2).$$

- For $\mathbf{x} = \mathbf{x}^*$, also $\mathbf{f}(\mathbf{x}) = \mathbf{0}$.
- Neglect nonlinear term and define method by

$$\mathbf{0} = \mathbf{f}(\mathbf{x}_k) + J(\mathbf{x}_k)(\mathbf{x}_{k+1} - \mathbf{x}_k).$$

- This is conceptually identical to the procedure (Chapter 3) for one function in one variable.
- **Algorithm:** Given an initial guess \mathbf{x}_0 ;
for $k = 0, 1, \dots$, until convergence
 solve $J(\mathbf{x}_k)\mathbf{p} = -\mathbf{f}(\mathbf{x}_k)$,
 set $\mathbf{x}_{k+1} = \mathbf{x}_k + \mathbf{p}$.
end

Newton's method

- **Derivation:** By Taylor series,

$$\mathbf{f}(\mathbf{x}) = \mathbf{f}(\mathbf{x}_k) + J(\mathbf{x}_k)(\mathbf{x} - \mathbf{x}_k) + \mathcal{O}(\|\mathbf{x} - \mathbf{x}_k\|^2).$$

- For $\mathbf{x} = \mathbf{x}^*$, also $\mathbf{f}(\mathbf{x}) = \mathbf{0}$.
- Neglect nonlinear term and define method by

$$\mathbf{0} = \mathbf{f}(\mathbf{x}_k) + J(\mathbf{x}_k)(\mathbf{x}_{k+1} - \mathbf{x}_k).$$

- This is conceptually identical to the procedure (Chapter 3) for one function in one variable.
- **Algorithm:** Given an initial guess \mathbf{x}_0 ;
for $k = 0, 1, \dots$, until convergence
 solve $J(\mathbf{x}_k)\mathbf{p} = -\mathbf{f}(\mathbf{x}_k)$,
 set $\mathbf{x}_{k+1} = \mathbf{x}_k + \mathbf{p}$.
end

Newton's method

- **Derivation:** By Taylor series,

$$\mathbf{f}(\mathbf{x}) = \mathbf{f}(\mathbf{x}_k) + J(\mathbf{x}_k)(\mathbf{x} - \mathbf{x}_k) + \mathcal{O}(\|\mathbf{x} - \mathbf{x}_k\|^2).$$


- For $\mathbf{x} = \mathbf{x}^*$, also $\mathbf{f}(\mathbf{x}) = \mathbf{0}$.
- Neglect nonlinear term and define method by

$$\mathbf{0} = \mathbf{f}(\mathbf{x}_k) + J(\mathbf{x}_k)(\mathbf{x}_{k+1} - \mathbf{x}_k).$$

- This is conceptually identical to the procedure (Chapter 3) for one function in one variable.
- **Algorithm:** Given an initial guess \mathbf{x}_0 ;
for $k = 0, 1, \dots$, until convergence
 solve $J(\mathbf{x}_k)\mathbf{p} = -\mathbf{f}(\mathbf{x}_k)$,
 set $\mathbf{x}_{k+1} = \mathbf{x}_k + \mathbf{p}$.
end

Example: a parabola meets a circle

$$\mathbf{x} = \begin{pmatrix} x_1 \\ x_2 \end{pmatrix}, \quad \mathbf{f}(\mathbf{x}) = \begin{pmatrix} x_1^2 - 2x_1 - x_2 + 1 \\ x_1^2 + x_2^2 - 1 \end{pmatrix}.$$




Two solutions: $(0,1)^T$ and $(1,0)^T$.

$$J(\mathbf{x}) = \begin{pmatrix} 2x_1 - 2 & -1 \\ 2x_1 & 2x_2 \end{pmatrix}$$

Example: a parabola meets a circle

$$\mathbf{x} = \begin{pmatrix} x_1 \\ x_2 \end{pmatrix}, \quad \mathbf{f}(\mathbf{x}) = \begin{pmatrix} x_1^2 - 2x_1 - x_2 + 1 \\ x_1^2 + x_2^2 - 1 \end{pmatrix}.$$



Two solutions: $(0, 1)^T$ and $(1, 0)^T$.

$$J(\mathbf{x}) = \begin{pmatrix} 2x_1 - 2 & -1 \\ 2x_1 & 2x_2 \end{pmatrix}$$

Example (cont.)

Using Newton's method for finding the two roots, varying the starting point \mathbf{x}_0 .

Stopping tolerance $\text{tol} = 1.\text{e-}7$.

- 1 Starting at $\mathbf{x}_0 = (0, 0)^T$ is bad because $J(\mathbf{x}_0)$ is singular!
- 2 Starting at $\mathbf{x}_0 = (1, 1)^T$ obtain root $(0, 1)^T$ in 5 iterations. Observe quadratic convergence.
- 3 Starting at $\mathbf{x}_0 = (-1, 1)^T$ obtain root $(1, 0)^T$ in 5 iterations. Observe quadratic convergence.

Example: two-point boundary value ordinary differential equation

- Consider the differential problem

$$\begin{aligned}u''(t) + e^{u(t)} &= 0, \quad 0 < t < 1, \\ u(0) &= u(1) = 0.\end{aligned}$$

- Discretize on a uniform mesh (grid) $t_i = ih$, $i = 0, 1, \dots, n+1$, where $(n+1)h = 1$:

$$\begin{aligned}\frac{v_{i+1} - 2v_i + v_{i-1}}{h^2} + e^{v_i} &= 0, \quad i = 1, 2, \dots, n. \\ v_0 &= v_{n+1} = 0.\end{aligned}$$

- This is a system of nonlinear equations, with $\mathbf{x} \leftarrow \mathbf{v}$ and $f_i(\mathbf{v}) = \frac{v_{i+1} - 2v_i + v_{i-1}}{h^2} + e^{v_i}$.

Example: two-point boundary value ordinary differential equation

- Consider the differential problem

$$\begin{aligned}u''(t) + e^{u(t)} &= 0, \quad 0 < t < 1, \\ u(0) &= u(1) = 0.\end{aligned}$$

- Discretize on a uniform mesh (grid) $t_i = ih$, $i = 0, 1, \dots, n+1$, where $(n+1)h = 1$:

$$\begin{aligned}\frac{v_{i+1} - 2v_i + v_{i-1}}{h^2} + e^{v_i} &= 0, \quad i = 1, 2, \dots, n. \\ v_0 &= v_{n+1} = 0.\end{aligned}$$

- This is a system of nonlinear equations, with $\mathbf{x} \leftarrow \mathbf{v}$ and $f_i(\mathbf{v}) = \frac{v_{i+1} - 2v_i + v_{i-1}}{h^2} + e^{v_i}$.


Newton for TPBVP

- Jacobian matrix is possibly large but tridiagonal

$$J = \frac{1}{h^2} \begin{pmatrix} -2 + h^2 e^{v_1} & 1 & & & & \\ 1 & -2 + h^2 e^{v_2} & 1 & & & \\ & \ddots & \ddots & \ddots & & \\ & & 1 & -2 + h^2 e^{v_{n-1}} & 1 & \\ & & & 1 & -2 + h^2 e^{v_n} \end{pmatrix}.$$

- Initial guess $\mathbf{v}_0 = \alpha (t_1(1 - t_1), \dots, t_n(1 - t_n))^T$.
- Take various values of α and see what happens, setting $\text{tol} = 1.e-8$, $n = 24$:
 - $\alpha = 0 \Rightarrow$ converges in 4 iterations
 - $\alpha = 10 \Rightarrow$ converges in 6 iterations
 - $\alpha = 20 \Rightarrow$ converges in 6 iterations to another solution
 - $\alpha = 50 \Rightarrow$ diverges

Two solutions



Outline

- Systems of nonlinear equations
- Unconstrained optimization
- *Constrained optimization

Multivariate Taylor expansion

- Consider the problem

$$\min_{\mathbf{x} \in \mathbb{R}^n} \phi(\mathbf{x})$$

Assume that $\phi(\mathbf{x})$ is smooth enough.

- Define **gradient** vector $\nabla \phi(\mathbf{x})$ and **Hessian** matrix $\nabla^2 \phi(\mathbf{x})$ by

$$\nabla \phi(\mathbf{x}) = \begin{pmatrix} \frac{\partial \phi}{\partial x_1} \\ \frac{\partial \phi}{\partial x_2} \\ \vdots \\ \frac{\partial \phi}{\partial x_n} \end{pmatrix}, \quad \nabla^2 \phi(\mathbf{x}) = \begin{pmatrix} \frac{\partial^2 \phi}{\partial x_1^2} & \frac{\partial^2 \phi}{\partial x_1 \partial x_2} & \cdots & \frac{\partial^2 \phi}{\partial x_1 \partial x_n} \\ \frac{\partial^2 \phi}{\partial x_2 \partial x_1} & \frac{\partial^2 \phi}{\partial x_2^2} & \cdots & \frac{\partial^2 \phi}{\partial x_2 \partial x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 \phi}{\partial x_n \partial x_1} & \frac{\partial^2 \phi}{\partial x_n \partial x_2} & \cdots & \frac{\partial^2 \phi}{\partial x_n^2} \end{pmatrix}.$$

- Taylor expansion near a point $\mathbf{x} \in \mathbb{R}^n$:

$$\phi(\mathbf{x} + \mathbf{p}) = \phi(\mathbf{x}) + \nabla \phi(\mathbf{x})^T \mathbf{p} + \frac{1}{2} \mathbf{p}^T \nabla^2 \phi(\mathbf{x}) \mathbf{p} + \mathcal{O}(\|\mathbf{p}\|^3).$$

Example

- Given the function

$$\phi(x_1, x_2) = x_1^4 - 2x_1^3x_2^2 + 4x_1x_2^3,$$

- the gradient at a point \mathbf{x} is

$$\nabla\phi(\mathbf{x}) = \begin{pmatrix} 4x_1^3 - 6x_1^2x_2^2 + 4x_2^3 \\ -4x_1^3x_2 + 12x_1x_2^2 \end{pmatrix};$$

- the Hessian matrix is

$$H(\mathbf{x}) = \nabla^2\phi(\mathbf{x}) = \begin{pmatrix} 12x_1^2 - 12x_1x_2^2 & -12x_1^2x_2 + 12x_2^2 \\ -12x_1^2x_2 + 12x_2^2 & -4x_1^3 + 24x_1x_2 \end{pmatrix}.$$


Critical points

- Taylor expansion near a suspected **minimum** point \mathbf{x}^* : in any direction \mathbf{p} with magnitude $\|\mathbf{p}\|$ small enough, must have $\phi(\mathbf{x}^* + \mathbf{p}) \geq \phi(\mathbf{x}^*)$.
- Hence

$$\nabla \phi(\mathbf{x}^*) = \mathbf{0}.$$

This defines a **critical point**.

- Similar condition also for a **maximum** or a **saddle** point.



Conditions for unconstrained minimum

$$\min_{\mathbf{x} \in \mathbb{R}^n} \phi(\mathbf{x}).$$

Assume that $\phi(\mathbf{x})$ is smooth enough.

- A necessary condition for having a **local minimum** at \mathbf{x}^* is that \mathbf{x}^* be a **critical point**

$$\nabla \phi(\mathbf{x}^*) = \mathbf{0},$$

and that the symmetric Hessian matrix $\nabla^2 \phi(\mathbf{x}^*)$ be **positive semi-definite**.

- A sufficient condition is that also $\nabla^2 \phi(\mathbf{x}^*)$ be **positive definite**.

Descent direction

- At point \mathbf{x} the vector \mathbf{p} is a descent direction if

$$\nabla\phi(\mathbf{x})^T \mathbf{p} = \sum_{i=1}^n p_i \frac{\partial\phi}{\partial x_i} < 0.$$

- A small step in a descent direction gives reduction in the objective function:

$$\phi(\mathbf{x} + \alpha\mathbf{p}) < \phi(\mathbf{x})$$

for scalar $0 < \alpha \ll 1$.

- Therefore, we can construct an iterative method that keeps reducing ϕ until convergence by using descent directions and controlled step sizes.

Descent direction

- At point \mathbf{x} the vector \mathbf{p} is a descent direction if

$$\nabla\phi(\mathbf{x})^T \mathbf{p} = \sum_{i=1}^n p_i \frac{\partial\phi}{\partial x_i} < 0.$$

- A small step in a descent direction gives reduction in the objective function:

$$\phi(\mathbf{x} + \alpha\mathbf{p}) < \phi(\mathbf{x})$$

for scalar $0 < \alpha \ll 1$.

- Therefore, we can construct an iterative method that keeps reducing ϕ until convergence by using descent directions and controlled step sizes.

Gradient-based methods

Consider

$$\begin{aligned}\mathbf{x}_{k+1} &= \mathbf{x}_k + \alpha_k \mathbf{p}_k, \quad \text{where} \\ \mathbf{p}_k &= -B_k^{-1} \nabla \phi(\mathbf{x}_k).\end{aligned}$$

If B_k is symmetric positive definite then \mathbf{p}_k is a descent direction. Note $\alpha_k > 0$.

- 1 Gradient descent: $B_k = I$. How to choose α_k ?
- 2 Nonlinear conjugate gradients.
- 3 Newton: $B_k = \nabla^2 \phi(\mathbf{x}_k)$. Set $\alpha_k = 1$ or damped Newton: search for $\alpha_k \leq 1$ guaranteeing descent.
- 4 Secant, or quasi-Newton.
- 5 Inexact Newton, Newton-CG.
- 6 Gauss-Newton for nonlinear least squares data fitting.

Gradient-based methods

Consider

$$\begin{aligned}\mathbf{x}_{k+1} &= \mathbf{x}_k + \alpha_k \mathbf{p}_k, \quad \text{where} \\ \mathbf{p}_k &= -B_k^{-1} \nabla \phi(\mathbf{x}_k).\end{aligned}$$

If B_k is symmetric positive definite then \mathbf{p}_k is a descent direction. Note $\alpha_k > 0$.

- ① Gradient descent: $B_k = I$. How to choose α_k ?
- ② Nonlinear conjugate gradients.
- ③ Newton: $B_k = \nabla^2 \phi(\mathbf{x}_k)$. Set $\alpha_k = 1$ or damped Newton: search for $\alpha_k \leq 1$ guaranteeing descent.
- ④ Secant, or quasi-Newton.
- ⑤ Inexact Newton, Newton-CG.
- ⑥ Gauss-Newton for nonlinear least squares data fitting.

Gradient-based methods

Consider

$$\begin{aligned}\mathbf{x}_{k+1} &= \mathbf{x}_k + \alpha_k \mathbf{p}_k, \quad \text{where} \\ \mathbf{p}_k &= -B_k^{-1} \nabla \phi(\mathbf{x}_k).\end{aligned}$$

If B_k is symmetric positive definite then \mathbf{p}_k is a descent direction. Note $\alpha_k > 0$.

- ① Gradient descent: $B_k = I$. How to choose α_k ?
- ② Nonlinear conjugate gradients.
- ③ Newton: $B_k = \nabla^2 \phi(\mathbf{x}_k)$. Set $\alpha_k = 1$ or damped Newton: search for $\alpha_k \leq 1$ guaranteeing descent.
- ④ Secant, or quasi-Newton.
- ⑤ Inexact Newton, Newton-CG.
- ⑥ Gauss-Newton for nonlinear least squares data fitting.

Gradient-based methods

Consider

$$\begin{aligned}\mathbf{x}_{k+1} &= \mathbf{x}_k + \alpha_k \mathbf{p}_k, \quad \text{where} \\ \mathbf{p}_k &= -B_k^{-1} \nabla \phi(\mathbf{x}_k).\end{aligned}$$

If B_k is symmetric positive definite then \mathbf{p}_k is a descent direction. Note $\alpha_k > 0$.

- ① Gradient descent: $B_k = I$. How to choose α_k ?
- ② Nonlinear conjugate gradients.
- ③ Newton: $B_k = \nabla^2 \phi(\mathbf{x}_k)$. Set $\alpha_k = 1$ or damped Newton: search for $\alpha_k \leq 1$ guaranteeing descent.
- ④ Secant, or quasi-Newton.
- ⑤ Inexact Newton, Newton-CG.
- ⑥ Gauss-Newton for nonlinear least squares data fitting.

Gradient-based methods

Consider

$$\begin{aligned}\mathbf{x}_{k+1} &= \mathbf{x}_k + \alpha_k \mathbf{p}_k, \quad \text{where} \\ \mathbf{p}_k &= -B_k^{-1} \nabla \phi(\mathbf{x}_k).\end{aligned}$$

If B_k is symmetric positive definite then \mathbf{p}_k is a descent direction. Note $\alpha_k > 0$.

- ① Gradient descent: $B_k = I$. How to choose α_k ?
- ② Nonlinear conjugate gradients.
- ③ Newton: $B_k = \nabla^2 \phi(\mathbf{x}_k)$. Set $\alpha_k = 1$ or damped Newton: search for $\alpha_k \leq 1$ guaranteeing descent.
- ④ Secant, or quasi-Newton.
- ⑤ Inexact Newton, Newton-CG.
- ⑥ Gauss-Newton for nonlinear least squares data fitting.

Newton's method

- Simply solve the nonlinear equations

$$\mathbf{f}(\mathbf{x}) \equiv \nabla \phi(\mathbf{x}) = \mathbf{0}$$

using Newton's method.

- At iteration k do

solve $\nabla^2 \phi(\mathbf{x}_k) \mathbf{p}_k = -\nabla \phi(\mathbf{x}_k)$ for \mathbf{p}_k

set $\mathbf{x}_{k+1} = \mathbf{x}_k + \mathbf{p}_k$.

- Note that if the Hessian matrix $\nabla^2 \phi(\mathbf{x}_k)$ is positive definite then \mathbf{p}_k minimizes a quadratic approximation to ϕ at \mathbf{x}_k :

$$\phi(\mathbf{x}_k + \mathbf{p}) \approx \phi(\mathbf{x}_k) + \nabla \phi(\mathbf{x}_k)^T \mathbf{p} + \frac{1}{2} \mathbf{p}^T \nabla^2 \phi(\mathbf{x}_k) \mathbf{p}.$$

- Furthermore, Newton's method has the following advantages:
 - locally, it converges quadratically;
 - (tautologically, it) retains Hessian sparsity.

Newton's method

- Simply solve the nonlinear equations

$$\mathbf{f}(\mathbf{x}) \equiv \nabla \phi(\mathbf{x}) = \mathbf{0}$$

using Newton's method.

- At iteration k do

solve $\nabla^2 \phi(\mathbf{x}_k) \mathbf{p}_k = -\nabla \phi(\mathbf{x}_k)$ for \mathbf{p}_k

set $\mathbf{x}_{k+1} = \mathbf{x}_k + \mathbf{p}_k$.

- Note that if the Hessian matrix $\nabla^2 \phi(\mathbf{x}_k)$ is positive definite then \mathbf{p}_k minimizes a quadratic approximation to ϕ at \mathbf{x}_k :

$$\phi(\mathbf{x}_k + \mathbf{p}) \approx \phi(\mathbf{x}_k) + \nabla \phi(\mathbf{x}_k)^T \mathbf{p} + \frac{1}{2} \mathbf{p}^T \nabla^2 \phi(\mathbf{x}_k) \mathbf{p}.$$

- Furthermore, Newton's method has the following advantages:
 - locally, it converges quadratically;
 - (tautologically, it) retains Hessian sparsity.

Newton's method deficiencies

- ❶ Requires existence of the Hessian.
- ❷ Requires evaluation of the Hessian.
- ❸ Requires solving a linear system at each iteration.
- ❹ $B_k = \nabla^2 \phi(\mathbf{x}_k)$ is symmetric but may not be positive definite.
- ❺ No control over convergence (does it converge? to a minimum point?)

Advanced methods are based on capitalizing on the strengths of Newton's method, while trying to weaken or eliminate its deficiencies.

(Weak) line search

- Suppose that \mathbf{p}_k is a descent direction at \mathbf{x}_k . Then for $\alpha_k > 0$ small enough, $\phi(\mathbf{x}_k + \alpha_k \mathbf{p}_k) < \phi(\mathbf{x}_k)$.
- Here is a simple algorithm (Armijo) for determining step size α_k , given descent direction \mathbf{p}_k :

Starting from $\alpha = \alpha_{max}$, repeat until sufficient decrease in ϕ is obtained,

$$\alpha \leftarrow \alpha/2.$$

- The result is $\alpha_k = \alpha$, and set $\mathbf{x}_{k+1} = \mathbf{x}_k + \alpha_k \mathbf{p}_k$.
- For Newton's method, set $\alpha_{max} = 1$ (and ensure that \mathbf{p}_k is a descent direction).
- For gradient descent, $B_k = I$ so \mathbf{p}_k is always a descent direction, but there is no obvious default value α_{max} .

Combining Newton and gradient descent

- Want the Newton efficiency. But what if $\nabla^2\phi(\mathbf{x}_k)$ is not positive definite?! – Need to ensure a descent search direction.
- So, consider mixing Newton and gradient descent:

$$B_k = \nabla^2\phi(\mathbf{x}_k) + \mu_k I.$$

For $\mu_k > 0$ large enough, this shifts the eigenvalues of B_k into positivity.

- Big question: choose scalar $\mu_k \geq 0$ adaptively, not too small and not too large; this leads to **trust region** methods.

Secant (quasi-Newton)

- Do not want to form or evaluate Hessian matrix explicitly; and *also*,
- want a positive definite B_k that is *easy to invert*.
- Note by Taylor's expansion

$$\nabla \phi(\mathbf{x}_k) \approx \nabla \phi(\mathbf{x}_{k+1}) - \nabla^2 \phi(\mathbf{x}_{k+1})(\mathbf{x}_{k+1} - \mathbf{x}_k).$$

The essential action of the Hessian is therefore in the direction of $\mathbf{w}_k = \mathbf{x}_{k+1} - \mathbf{x}_k$.

- Thus, equire

$$B_{k+1} \mathbf{w}_k = \mathbf{y}_k, \quad \mathbf{y}_k = \nabla \phi(\mathbf{x}_{k+1}) - \nabla \phi(\mathbf{x}_k).$$

- Obtain B_{k+1} as a positive definite rank-2 update of B_k , thus satisfying the above requirements.

Secant (quasi-Newton)

- Do not want to form or evaluate Hessian matrix explicitly; and *also*,
- want a positive definite B_k that is *easy to invert*.
- Note by Taylor's expansion

$$\nabla \phi(\mathbf{x}_k) \approx \nabla \phi(\mathbf{x}_{k+1}) - \nabla^2 \phi(\mathbf{x}_{k+1})(\mathbf{x}_{k+1} - \mathbf{x}_k).$$

The essential action of the Hessian is therefore in the direction of

$$\mathbf{w}_k = \mathbf{x}_{k+1} - \mathbf{x}_k.$$

- Thus, equire

$$B_{k+1} \mathbf{w}_k = \mathbf{y}_k, \quad \mathbf{y}_k = \nabla \phi(\mathbf{x}_{k+1}) - \nabla \phi(\mathbf{x}_k).$$

- Obtain B_{k+1} as a positive definite rank-2 update of B_k , thus satisfying the above requirements.

Secant (quasi-Newton)

- Do not want to form or evaluate Hessian matrix explicitly; and *also*,
- want a positive definite B_k that is *easy to invert*.
- Note by Taylor's expansion

$$\nabla \phi(\mathbf{x}_k) \approx \nabla \phi(\mathbf{x}_{k+1}) - \nabla^2 \phi(\mathbf{x}_{k+1})(\mathbf{x}_{k+1} - \mathbf{x}_k).$$

The essential action of the Hessian is therefore in the direction of

$$\mathbf{w}_k = \mathbf{x}_{k+1} - \mathbf{x}_k.$$

- Thus, equire

$$B_{k+1} \mathbf{w}_k = \mathbf{y}_k, \quad \mathbf{y}_k = \nabla \phi(\mathbf{x}_{k+1}) - \nabla \phi(\mathbf{x}_k).$$

- Obtain B_{k+1} as a positive definite rank-2 update of B_k , thus satisfying the above requirements.

BFGS method

Update $G_k = B_k^{-1}$ directly.

Choose \mathbf{x}_0 and G_0 (e.g., $G_0 = I$)

for $k = 0, 1, \dots$, until convergence

$$\mathbf{p}_k = -G_k \nabla \phi(\mathbf{x}_k)$$

find a suitable step size α_k

$$\mathbf{x}_{k+1} = \mathbf{x}_k + \alpha_k \mathbf{p}_k$$

$$\mathbf{w}_k = \alpha_k \mathbf{p}_k$$

$$\mathbf{y}_k = \nabla \phi(\mathbf{x}_{k+1}) - \nabla \phi(\mathbf{x}_k)$$

$$G_{k+1} = \left(I - \frac{\mathbf{w}_k \mathbf{y}_k^T}{\mathbf{y}_k^T \mathbf{w}_k}\right) G_k \left(I - \frac{\mathbf{y}_k \mathbf{w}_k^T}{\mathbf{y}_k^T \mathbf{w}_k}\right) + \frac{\mathbf{w}_k \mathbf{w}_k^T}{\mathbf{y}_k^T \mathbf{w}_k}.$$

BFGS method

Update $G_k = B_k^{-1}$ directly.

Choose \mathbf{x}_0 and G_0 (e.g., $G_0 = I$)

for $k = 0, 1, \dots$, until convergence

$$\mathbf{p}_k = -G_k \nabla_{\phi}(\mathbf{x}_k)$$

find a suitable step size α_k

$$\mathbf{x}_{k+1} = \mathbf{x}_k + \alpha_k \mathbf{p}_k$$

$$\mathbf{w}_k = \alpha_k \mathbf{p}_k$$

$$\mathbf{y}_k = \nabla_{\phi}(\mathbf{x}_{k+1}) - \nabla_{\phi}(\mathbf{x}_k)$$

$$G_{k+1} = \left(I - \frac{\mathbf{w}_k \mathbf{y}_k^T}{\mathbf{y}_k^T \mathbf{w}_k}\right) G_k \left(I - \frac{\mathbf{y}_k \mathbf{w}_k^T}{\mathbf{y}_k^T \mathbf{w}_k}\right) + \frac{\mathbf{w}_k \mathbf{w}_k^T}{\mathbf{y}_k^T \mathbf{w}_k}.$$

Secant (quasi-Newton) pros and cons

- Cheap update of $G_k = B_k^{-1}$ using rank-2 updates.
- Local superlinear convergence.
- Only descent directions because $G_k = B_k^{-1}$ symmetric positive definite.
- The method of choice for most problems.
- May lose Hessian sparsity, though.
- Limited memory versions L-BFGS exist for large, sparse matrices. They have their pros and cons...

Secant (quasi-Newton) pros and cons

- Cheap update of $G_k = B_k^{-1}$ using rank-2 updates.
- Local superlinear convergence.
- Only descent directions because $G_k = B_k^{-1}$ symmetric positive definite.
- The method of choice for most problems.
- May lose Hessian sparsity, though.
- Limited memory versions L-BFGS exist for large, sparse matrices. They have their pros and cons...

Inexact Newton, Newton-CG

- If the problem is large and has a sparse Hessian, may want to use Newton.
But how to solve the linear system at each iteration?
- Use iterative method (Chapter 7) for the linear system at each iteration:
 apply preconditioned conjugate gradient iterations towards solving
$$\nabla^2 \phi(\mathbf{x}_k) \mathbf{p}_k = -\nabla \phi(\mathbf{x}_k) \text{ for } \mathbf{p}_k$$
set $\mathbf{x}_{k+1} = \mathbf{x}_k + \mathbf{p}_k$.
- Newton - outer iteration
 PCG - inner iteration
- How many inner iterations?! (intuitively, fewer when far, more when close) –
 leads to inexact Newton.

Nonlinear least squares

$$\min_{\mathbf{x}} \phi(\mathbf{x}) = \frac{1}{2} \|\mathbf{g}(\mathbf{x}) - \mathbf{b}\|^2,$$

where \mathbf{b} is data (m values) and \mathbf{g} a nonlinear function of n parameters \mathbf{x} .

$$A(\mathbf{x}) = \begin{pmatrix} \frac{\partial g_1}{\partial x_1} & \frac{\partial g_1}{\partial x_2} & \cdots & \frac{\partial g_1}{\partial x_n} \\ \frac{\partial g_2}{\partial x_1} & \frac{\partial g_2}{\partial x_2} & \cdots & \frac{\partial g_2}{\partial x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial g_{m-1}}{\partial x_1} & \frac{\partial g_{m-1}}{\partial x_2} & \cdots & \frac{\partial g_{m-1}}{\partial x_n} \\ \frac{\partial g_m}{\partial x_1} & \frac{\partial g_m}{\partial x_2} & \cdots & \frac{\partial g_m}{\partial x_n} \end{pmatrix},$$

Then

$$\nabla \phi(\mathbf{x}) = A(\mathbf{x})^T (\mathbf{g}(\mathbf{x}) - \mathbf{b}).$$

Nonlinear least squares

$$\min_{\mathbf{x}} \phi(\mathbf{x}) = \frac{1}{2} \|\mathbf{g}(\mathbf{x}) - \mathbf{b}\|^2,$$

where \mathbf{b} is data (m values) and \mathbf{g} a nonlinear function of n parameters \mathbf{x} .

$$A(\mathbf{x}) = \begin{pmatrix} \frac{\partial g_1}{\partial x_1} & \frac{\partial g_1}{\partial x_2} & \cdots & \frac{\partial g_1}{\partial x_n} \\ \frac{\partial g_2}{\partial x_1} & \frac{\partial g_2}{\partial x_2} & \cdots & \frac{\partial g_2}{\partial x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial g_{m-1}}{\partial x_1} & \frac{\partial g_{m-1}}{\partial x_2} & \cdots & \frac{\partial g_{m-1}}{\partial x_n} \\ \frac{\partial g_m}{\partial x_1} & \frac{\partial g_m}{\partial x_2} & \cdots & \frac{\partial g_m}{\partial x_n} \end{pmatrix},$$

Then

$$\nabla \phi(\mathbf{x}) = A(\mathbf{x})^T (\mathbf{g}(\mathbf{x}) - \mathbf{b}).$$

Gauss-Newton

Hessian matrix

$$\nabla^2 \phi(\mathbf{x}) = A(\mathbf{x})^T A(\mathbf{x}) + L(\mathbf{x}),$$

where L is $n \times n$

$$L_{i,j} = \sum_{l=1}^m \frac{\partial^2 g_l}{\partial x_i \partial x_j} (g_l - b_l).$$

L can be ugly. Appears in Newton's method!

Gauss-Newton: drop the ugly term. Define iteration by

$$\begin{aligned} [A(\mathbf{x}_k)^T A(\mathbf{x}_k)] \mathbf{p}_k &= -\nabla \phi(\mathbf{x}_k) \\ \mathbf{x}_{k+1} &= \mathbf{x}_k + \mathbf{p}_k. \end{aligned}$$

These are normal equations for

$$\min_{\mathbf{p}} \|A(\mathbf{x}_k)\mathbf{p} - (\mathbf{b} - \mathbf{g}(\mathbf{x}_k))\|,$$

so at each iteration solve a linear least squares problem (Chapter 6).

Gauss-Newton

Hessian matrix

$$\nabla^2 \phi(\mathbf{x}) = A(\mathbf{x})^T A(\mathbf{x}) + L(\mathbf{x}),$$

where L is $n \times n$

$$L_{i,j} = \sum_{l=1}^m \frac{\partial^2 g_l}{\partial x_i \partial x_j} (g_l - b_l).$$

L can be ugly. Appears in Newton's method!

Gauss-Newton: drop the ugly term. Define iteration by

$$\begin{aligned} [A(\mathbf{x}_k)^T A(\mathbf{x}_k)] \mathbf{p}_k &= -\nabla \phi(\mathbf{x}_k) \\ \mathbf{x}_{k+1} &= \mathbf{x}_k + \mathbf{p}_k. \end{aligned}$$

These are normal equations for

$$\min_{\mathbf{p}} \|A(\mathbf{x}_k)\mathbf{p} - (\mathbf{b} - \mathbf{g}(\mathbf{x}_k))\|,$$

so at each iteration solve a linear least squares problem (Chapter 6).

Gauss-Newton vs. Newton

- The Gauss-Newton direction, unlike Newton's, is guaranteed to be a descent direction. This is because $A^T A$ is positive definite even when $A^T A + L$ is not.
- The Gauss-Newton iteration is cheaper and can be better conditioned (more stable) than Newton's iteration.
- The convergence order of Gauss-Newton is only linear, as the difference between it and Newton's iteration does not vanish in the limit.
- Gauss-Newton converges faster for problems where the model fits the data well! This is because then $\|\mathbf{g}(\mathbf{x}) - \mathbf{b}\|$ is “small” near the solution, hence L is small and Gauss-Newton is closer to Newton.

Outline

- Systems of nonlinear equations
- Unconstrained optimization
- *Constrained optimization

Constrained problem

- General form

$$\min_{\mathbf{x} \in \Omega} \phi(\mathbf{x}), \quad \text{where}$$

$$\Omega = \{\mathbf{x} \in \mathbb{R}^n \mid c_i(\mathbf{x}) = 0, i \in \mathcal{E}, \quad c_i(\mathbf{x}) \geq 0, i \in \mathcal{I}\}.$$

- **Equality constraints**: reducing space; algebraic; domain Ω has empty interior.
- **Inequality constraints**: combinatorial; if \mathcal{E} empty then domain Ω can have nonempty interior.
- **Active set**

$$\mathcal{A}(\mathbf{x}) = \mathcal{E} \cup \{i \in \mathcal{I} \mid c_i(\mathbf{x}) = 0\}.$$

Consider problems where $\mathcal{A}(\mathbf{x}^*)$ is nonempty.

KKT conditions for a minimum

- Assume **constraint qualification**.

- Lagrangian**

$$\mathcal{L}(\mathbf{x}, \boldsymbol{\lambda}) = \phi(\mathbf{x}) - \sum_{i \in \mathcal{E} \cup \mathcal{I}} \lambda_i c_i(\mathbf{x}).$$

- KKT conditions** necessary for a minimum:

$$\begin{aligned}\nabla_{\mathbf{x}} \mathcal{L}(\mathbf{x}^*, \boldsymbol{\lambda}^*) &= \mathbf{0}, \\ c_i(\mathbf{x}^*) &= 0, \quad \forall i \in \mathcal{E}, \\ c_i(\mathbf{x}^*) &\geq 0, \quad \forall i \in \mathcal{I}, \\ \lambda_i^* &\geq 0, \quad \forall i \in \mathcal{I}, \\ \lambda_i^* c_i(\mathbf{x}^*) &= 0, \quad \forall i \in \mathcal{E} \cup \mathcal{I}.\end{aligned}$$

Active set methods

- Assuming solution is on $\partial\Omega$, search for the optimum along the boundary.
- For inequality constraints, keep track of $\mathcal{A}(\mathbf{x}_k)$, shuffling constraints in and out of the active set.
- e.g. quadratic programming (QP): quadratic objective function subject to linear inequality constraints.
- Sequential quadratic programming (SQP): At each iteration solve QP for search direction \mathbf{p}_k at $(\mathbf{x}_k, \lambda_k)$

$$\min_{\mathbf{p}} \quad \frac{1}{2} \mathbf{p}^T W_k \mathbf{p} + \nabla \phi(\mathbf{x}_k)^T \mathbf{p},$$

$$c_i(\mathbf{x}_k) + \nabla c_i(\mathbf{x}_k)^T \mathbf{p} = 0, \quad i \in \mathcal{E}, \quad c_i(\mathbf{x}_k) + \nabla c_i(\mathbf{x}_k)^T \mathbf{p} \geq 0, \quad i \in \mathcal{I}.$$

Active set methods

- Assuming solution is on $\partial\Omega$, search for the optimum along the boundary.
- For inequality constraints, keep track of $\mathcal{A}(\mathbf{x}_k)$, shuffling constraints in and out of the active set.
- e.g. **quadratic programming (QP)**: quadratic objective function subject to linear inequality constraints.
- **Sequential quadratic programming (SQP)**: At each iteration solve QP for search direction \mathbf{p}_k at $(\mathbf{x}_k, \lambda_k)$

$$\min_{\mathbf{p}} \quad \frac{1}{2} \mathbf{p}^T W_k \mathbf{p} + \nabla \phi(\mathbf{x}_k)^T \mathbf{p},$$

$$c_i(\mathbf{x}_k) + \nabla c_i(\mathbf{x}_k)^T \mathbf{p} = 0, \quad i \in \mathcal{E}, \quad c_i(\mathbf{x}_k) + \nabla c_i(\mathbf{x}_k)^T \mathbf{p} \geq 0, \quad i \in \mathcal{I}.$$

Active set methods

- Assuming solution is on $\partial\Omega$, search for the optimum along the boundary.
- For inequality constraints, keep track of $\mathcal{A}(\mathbf{x}_k)$, shuffling constraints in and out of the active set.
- e.g. **quadratic programming (QP)**: quadratic objective function subject to linear inequality constraints.
- **Sequential quadratic programming (SQP)**: At each iteration solve QP for search direction \mathbf{p}_k at $(\mathbf{x}_k, \boldsymbol{\lambda}_k)$

$$\min_{\mathbf{p}} \quad \frac{1}{2} \mathbf{p}^T W_k \mathbf{p} + \nabla_{\phi}(\mathbf{x}_k)^T \mathbf{p},$$

$$c_i(\mathbf{x}_k) + \nabla \mathbf{c}_i(\mathbf{x}_k)^T \mathbf{p} = 0, \quad i \in \mathcal{E}, \quad c_i(\mathbf{x}_k) + \nabla \mathbf{c}_i(\mathbf{x}_k)^T \mathbf{p} \geq 0, \quad i \in \mathcal{I}.$$

Interior point and other methods

- Penalty methods

$$\min_{\mathbf{x}} \psi(\mathbf{x}, \mu) = \phi(\mathbf{x}) + \frac{1}{2\mu} \sum_{i \in \mathcal{E}} c_i^2(\mathbf{x}),$$

where $\mu \downarrow 0$.

- Barrier methods

$$\min_{\mathbf{x}} \psi(\mathbf{x}, \mu) = \phi(\mathbf{x}) - \mu \sum_{i \in \mathcal{I}} \log c_i(\mathbf{x}),$$

where $\mu \downarrow 0$.

- Augmented Lagrangian

$$\min_{\mathbf{x}} \psi(\mathbf{x}, \lambda, \mu) = \phi(\mathbf{x}) - \sum_{i \in \mathcal{E}} \lambda_i c_i(\mathbf{x}) + \frac{1}{2\mu} \sum_{i \in \mathcal{E}} c_i^2(\mathbf{x}).$$

Given estimates λ_k, μ_k , solve the unconstrained minimization problem for $\mathbf{x} = \mathbf{x}_{k+1}$, then update the multipliers to λ_{k+1}, μ_{k+1} .

Interior point and other methods

- Penalty methods

$$\min_{\mathbf{x}} \psi(\mathbf{x}, \mu) = \phi(\mathbf{x}) + \frac{1}{2\mu} \sum_{i \in \mathcal{E}} c_i^2(\mathbf{x}),$$

where $\mu \downarrow 0$.

- Barrier methods

$$\min_{\mathbf{x}} \psi(\mathbf{x}, \mu) = \phi(\mathbf{x}) - \mu \sum_{i \in \mathcal{I}} \log c_i(\mathbf{x}),$$

where $\mu \downarrow 0$.

- Augmented Lagrangian

$$\min_{\mathbf{x}} \psi(\mathbf{x}, \lambda, \mu) = \phi(\mathbf{x}) - \sum_{i \in \mathcal{E}} \lambda_i c_i(\mathbf{x}) + \frac{1}{2\mu} \sum_{i \in \mathcal{E}} c_i^2(\mathbf{x}).$$

Given estimates λ_k, μ_k , solve the unconstrained minimization problem for $\mathbf{x} = \mathbf{x}_{k+1}$, then update the multipliers to λ_{k+1}, μ_{k+1} .

Interior point and other methods

- Penalty methods

$$\min_{\mathbf{x}} \psi(\mathbf{x}, \mu) = \phi(\mathbf{x}) + \frac{1}{2\mu} \sum_{i \in \mathcal{E}} c_i^2(\mathbf{x}),$$

where $\mu \downarrow 0$.

- Barrier methods

$$\min_{\mathbf{x}} \psi(\mathbf{x}, \mu) = \phi(\mathbf{x}) - \mu \sum_{i \in \mathcal{I}} \log c_i(\mathbf{x}),$$

where $\mu \downarrow 0$.

- Augmented Lagrangian

$$\min_{\mathbf{x}} \psi(\mathbf{x}, \boldsymbol{\lambda}, \mu) = \phi(\mathbf{x}) - \sum_{i \in \mathcal{E}} \lambda_i c_i(\mathbf{x}) + \frac{1}{2\mu} \sum_{i \in \mathcal{E}} c_i^2(\mathbf{x}).$$

Given estimates $\boldsymbol{\lambda}_k, \mu_k$, solve the unconstrained minimization problem for $\mathbf{x} = \mathbf{x}_{k+1}$, then update the multipliers to $\boldsymbol{\lambda}_{k+1}, \mu_{k+1}$.