

Computer Operating Systems

Practice Session 2: Booting Sequence and /proc File System

T. Tolga Sarı (sarita@itu.edu.tr)

Sultan Çoğay (cogay@itu.edu.tr)

Doğukan Arslan (arslan.dogukan@itu.edu.tr)

March 8, 2022

Today

Operating Systems, PS 2

PC Booting Sequence

Master Boot Record - MBR

Preloading Sectors

Linux /proc directory

When you press the power button...

- ▶ The system which starts the PC after the power button is pressed is called the boot loader (e.g. BIOS (**B**asic **I**nput **O**utput **S**ystem))
- ▶ BIOS is a series of information which is stored on hardware (ROM).

Initial processes

First checks and operations:

- ▶ Power Good Signal is the signal that is generated when the power supply reaches its required operating conditions (which is typically +5V).
- ▶ CPU is ready for operating. The First place to look up is the BIOS ROM for the start up program. Typically, the ROM ends with the memory space including the *jump* command.
- ▶ The First operation the BIOS performs is to check the system: a process called Power On Self Test (POST). The hardware is checked for any potential malfunction before the system starts.
- ▶ The graphics card is started via searching for its BIOS.

BIOS checks

ROMs of the remaining peripherals is searched for a BIOS.

- ▶ Typically, the BIOSes of the IDE/ATA hard drives are found and executed.
- ▶ If any other peripheral has a BIOS, then it is also executed similarly.

Startup screen

BIOS visualizes its startup screen. This startup screen contains the following information:

- ▶ BIOS producer and version number
- ▶ BIOS date
- ▶ Keys to enter the BIOS Setup
- ▶ System logo
- ▶ BIOS serial number
- ▶ <http://www.wimsbios.com/> (an online BIOS scan)

BIOS tests

- ▶ BIOS performs various tests on system such as memory count test.
- ▶ The user is informed on any errors encountered at this point.
- ▶ *"Keyboard not found, press F1 to continue..."*

Persistent system information

- ▶ After previous operations, BIOS reads the system date, system time and peripherals list from the CMOS memory on the mainboard.
- ▶ CMOS integrated circuits require very low power, thus they are able to store their memories for very extended periods with a standard battery. In PCs, CMOS integrated circuits are typically used for storing the data such as date and time, which need to be unaffected from power failures.
- ▶ By reading the information stored in the CMOS, the PC learns which hard drives are connected and in which order they should be checked for a proper startup sequence. Therefore, it is able to start the operating system properly.

Master Boot Record

- ▶ If the booting will be performed using a hard drive, Cylinder 0, Head 0, Sector 1 which is called *Master Boot Record* is read.
- ▶ At this point, BIOS is disengaged.
- ▶ In order to load the OS, system copies the first 512 bytes of the first hard drive into the memory and executes the code existing at the beginning of this section. Information included is related to the further booting operations. That is why it is called MBR.

PC Booting Sequence

Up to this point, booting operations are independent of the installed operating system and are the same for all PCs.

Master Boot Record - MBR

- ▶ The organization of the MBR has a very standard structure irrespective of the type of the installed operating system:
 - ▶ First part of 446 bytes are reserved for the program code.
 - ▶ Latter 64 bytes includes a partition table containing 4 partitions.
 - ▶ Last 2 bytes includes a special number (magic number AA55). An MBR having a different number is not validated by BIOS and any operating system.
- ▶ Program starts booting sequence by looking at the partition table and deciding which partition to be used for the startup. Then, program transfer the flow control to the specified partitions preloading sector (boot sector).

Locations of the preloading sectors

- ▶ Preloading sectors are the first sectors of the hard discs (a.k.a. boot sectors). They provide a space (512 bytes) for the code to start the operating system in that portion. Additionally, they include some basic information on the file system.
- ▶ A valid preloading sector (likewise in MBR) includes a special number stored in last 2 bytes (AA55).

Linux Boot Loaders

In Linux, different boot loaders can be written to different preloading sectors.

- ▶ LILO (Linux Loader) - GRUB (Grand Unified Boot Loader)
 - ▶ Are responsible for the loading of the system and conveying the control to the kernel.
 - ▶ Supports many operating systems and file systems.
- ▶ LILO (Linux Loader) - GRUB (Grand Unified Boot Loader) differences
 - ▶ LILO, does not provide interactive command interface like GRUB.
 - ▶ LILO does not support booting from network: GRUB does.
 - ▶ In LILO, with an erroneous modification in the config file, MBR with an improper configuration may cause the system to be un-bootable. In GRUB, on the occurrence of such condition, system passes to the interactive command interface.

Kernel functions and /proc

- ▶ Linux kernel has two basic functionalities:
 - ▶ Control the access to the hardware
 - ▶ Determine when and how the processes will interact with these entities
- ▶ /proc folder contains files about the current status of the kernel.
- ▶ Information about hardware and active processes can be retrieved from files under /proc directory.
- ▶ /proc folder is on the virtual file system.
- ▶ In virtual file systems, information is kept in memory: do not take any place in discs.
- ▶ In virtual file systems, files act and seem like usual files.

/proc directory contents

```

root@koknar:/proc - Kabuk - Konsole
Oturum Düzenle Görüntüle Yer İmleri Ayarlar Yardım

[root@koknar root]# cd /proc/
[root@koknar proc]# ls
1      1742  2142  2511  2713  3      crypto      irq          pci
10     1932  2147  2514  2714  4      devices     kcore        scsi
110    1947  2151  2515  2715  5      diskstats   kmsg         self
151    1966  2154  2645  2716  6      dma          loadavg      slabinfo
152    1975  2157  2688  2719  7      driver       locks        stat
153    2      2159  2689  2720  8      execdomains  mdstat       swaps
1551   2006  2160  2691  2754  9      fb           meminfo      sys
1555   2037  2362  2692  2787  acpi        filesystems  misc         sysrq-trigger
1576   2059  2373  2694  2863  asound      fs           modules      sysvipc
1596   2069  2403  2703  2864  buddyinfo  ide          mounts        tty
1623   2078  2468  2704  2893  bus         interrupts  mtrr          uptime
166    2097  2506  2710  2894  cmdline    iomem        net           version
1730   2113  2509  2712  2911  cpuinfo     ioports      partitions    vmstat

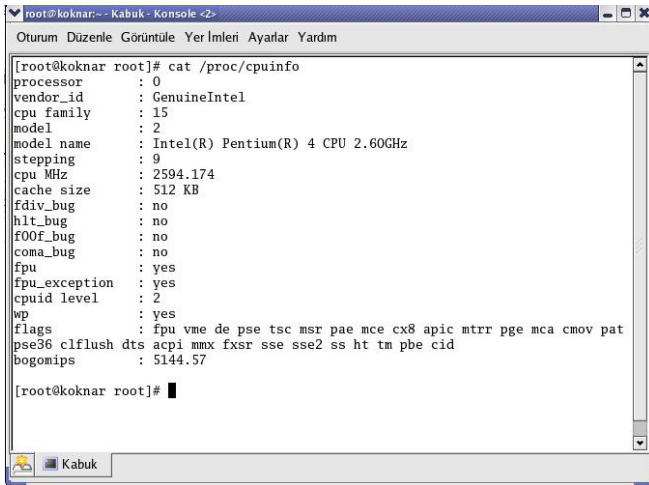
[root@koknar proc]#

```

Properties of the files under /proc

- ▶ Files under /proc folder are updated continuously. Therefore:
 - ▶ Most of them always have size of 0 bytes.
 - ▶ The date and settings for the last access records of most of them reflect the current date and time.
- ▶ Most of the files are accessible to only 'root'.
- ▶ Files under /proc folder include many information about the system. Such as:
 - ▶ uptime, version, kcore (displays a value given in bytes representing the size of the physical memory)...
 - ▶ `cat /proc/cpuinfo`

Accessing CPU information



The screenshot shows a terminal window titled "root@koknar:~ - Kabuk - Konsolle <2>". The window has a menu bar with "Oturum", "Düzenle", "Görüntüle", "Yer İmleri", "Ayarlar", and "Yardım". The terminal content shows the command `cat /proc/cpuinfo` being executed, displaying the following CPU information:

```
[root@koknar root]# cat /proc/cpuinfo
processor       : 0
vendor_id      : GenuineIntel
cpu family     : 15
model          : 2
model name     : Intel(R) Pentium(R) 4 CPU 2.60GHz
stepping      : 9
cpu MHz        : 2594.174
cache size     : 512 KB
fdiv_bug       : no
hlt_bug        : no
f00f_bug       : no
coma_bug       : no
fpu            : yes
fpu_exception  : yes
cpuid level    : 2
wp             : yes
flags          : fpu vme de pse tsc msr pae mce cx8 apic mtrr pge mca cmov pat
                 pse36 clflush dts acpi mmx fxsr sse sse2 ss ht tm pbe cid
bogomips       : 5144.57

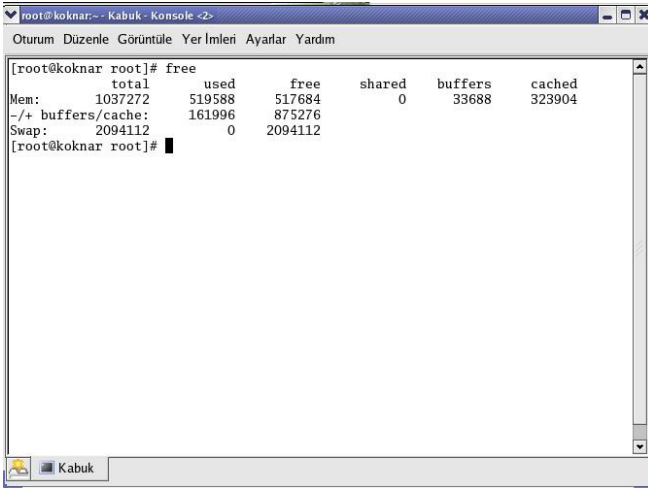
[root@koknar root]#
```

The terminal window has a taskbar at the bottom with a "Kabuk" icon.

Monitoring memory space

- ▶ Some files under /proc are hard to read with naked eye. Therefore, we use auxiliary commands:
- ▶ Example: `free` gives information about memory space:
 - ▶ Swap space
 - ▶ Free and used portions of the physical memory
 - ▶ Buffers and cache consumed by the kernel

free command



The screenshot shows a terminal window titled "root@koknar:~ - Kabuk - Konsole <2>". The window has a menu bar with "Oturum", "Düzenle", "Görüntüle", "Yer İmleri", "Ayarlar", and "Yardım". The terminal content shows the command "free" being executed, resulting in a table of memory statistics. The prompt is "[root@koknar root]#".

```
[root@koknar root]# free
```

	total	used	free	shared	buffers	cached
Mem:	1037272	519588	517684	0	33688	323904
-/+ buffers/cache:		161996	875276			
Swap:	2094112	0	2094112			

```
[root@koknar root]#
```

top command

```

root@koknar:~ - Kabuk - Konsole <2>
Oturum Düzenle Görüntüle Yer İmleri Ayarlar Yardım

top - 11:31:52 up 37 min, 3 users, load average: 0.61, 0.33, 0.21
Tasks: 71 total, 1 running, 70 sleeping, 0 stopped, 0 zombie
Cpu(s): 3.0% us, 0.7% sy, 0.0% ni, 96.3% id, 0.0% wa, 0.0% hi, 0.0% si
top - 11:32:11 up 37 min, 3 users, load average: 0.50, 0.32, 0.21
Tasks: 71 total, 2 running, 69 sleeping, 0 stopped, 0 zombie
top - 11:32:20 up 37 min, 3 users, load average: 0.42, 0.31, 0.21
Tasks: 71 total, 2 running, 69 sleeping, 0 stopped, 0 zombie
top - 11:34:57 up 40 min, 3 users, load average: 0.16, 0.22, 0.18
Tasks: 71 total, 1 running, 70 sleeping, 0 stopped, 0 zombie
Cpu(s): 2.7% us, 0.0% sy, 0.0% ni, 97.3% id, 0.0% wa, 0.0% hi, 0.0% si
Mem: 1037272k total, 552004k used, 485268k free, 33896k buffers
Swap: 2094112k total, 0k used, 2094112k free, 323916k cached

  PID USER      PR  NI  VIRT  RES  SHR  S %CPU  %MEM    TIME+  COMMAND
 2373 root        15   0 163m 15m 149m S  2.3  1.6   1:22.82 X
3043 root        15   0 29804 14m 26m S  0.3  1.4   0:00.78 kdeinit
   1 root        16   0 2984  464 1316 S  0.0  0.0   0:05.50 init
   2 root        34  19   0    0   0 S  0.0  0.0   0:00.00 ksoftirqd/0
   3 root        5 -10   0    0   0 S  0.0  0.0   0:00.01 events/0
   4 root        5 -10   0    0   0 S  0.0  0.0   0:00.00 kblockd/0
   6 root       12 -10   0    0   0 S  0.0  0.0   0:00.00 khelper
   5 root       15   0   0    0   0 S  0.0  0.0   0:00.00 khubd
   7 root       20   0   0    0   0 S  0.0  0.0   0:00.00 pdflush
   8 root       15   0   0    0   0 S  0.0  0.0   0:00.01 pdflush
  10 root       12 -10   0    0   0 S  0.0  0.0   0:00.00 aio/0

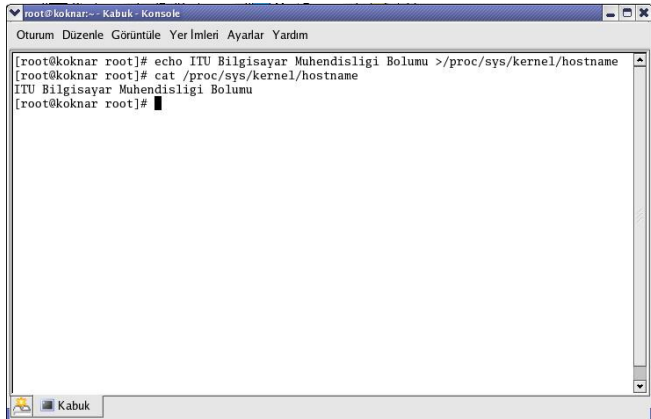
```

- ▶ PR: Priority level
- ▶ NI: Nice parameter, used in scheduling (Negative values - higher priority)
- ▶ VIRT: Virtual memory space used by the process
- ▶ SHR: How much virtual memory can be shared
- ▶ RES: Usage of the physical memory

Writing into the files under /proc

- ▶ Most of the time, these files are read-only.
- ▶ Some of them may be modified in order to configure some kernel parameters.
- ▶ Since the files are virtual, shell commands are needed for performing the modifications.

Writing to a file under /proc using echo command



```
root@koknar:~ - Kabuk - Konsole
Oturum Düzenle Görüntüle Yer İmleri Ayarlar Yardım

[root@koknar root]# echo ITU Bilgisayar Muhendisligi Bolunu >/proc/sys/kernel/hostname
[root@koknar root]# cat /proc/sys/kernel/hostname
ITU Bilgisayar Muhendisligi Bolunu
[root@koknar root]#
```

Process folders under /proc

- Each running process has a folder under /proc.

```

root@koknar:/proc/3339 - Kabuk - Konsole
Oturum Düzenle Görüntüle Yer İmleri Ayarlar Yardım

[root@koknar root]# cd /proc/
[root@koknar proc]# ls
1      1932 2147 3324 3353 6      diskstats  kmsg      self
10     1947 2151 3332 3355 6800    dma        loadavg   slabinfo
110    1966 2154 3335 3359 7      driver     locks     stat
151    1975 2157 3336 3360 8      execdomains mdstat    swaps
152    2    2159 3338 3378 9      fb         meminfo   sys
153    2006 2160 3339 3379  acpi      filesystems misc       sysrq-trigger
1551   2037 2362 3341 3384  asound    fs         modules   sysvipc
1555   2059 2373 3343 3387  buddyinfo ide         mounts     tty
1576   2069 3    3347 3432  bus        interrupts mtrr       uptime
1596   2078 3214 3349 3439  cmdline   iomem      net        version
1623   2097 3279 3350 3539  cpuinfo    ioports    partitions vmstat
166    2113 3318 3351 4      crypto     irq        pci
1730   2142 3323 3352 5      devices    kcore      scsi

[root@koknar proc]# cd 3339
[root@koknar 3339]# ls
attr  cmdline  environ  fd      mem      root  statm  task
auxv  cwd      exe      maps    mounts   stat  status wchan
[root@koknar 3339]#

```

References

- ▶ <http://www.redhat.com/docs/manuals/linux/RHL-9-Manual/ref-guide/ch-proc.html>
- ▶ <http://www.kernelnewbies.org/documents/kdoc/procfs-guide/lkprocfsguide.html>
- ▶ <http://www.redhat.com/docs/manuals/linux/RHL-9-Manual/ref-guide/s1-proc-topfiles.html>
- ▶ <http://www.belgeler.org/>